

```
!pip install bayesian-optimization
```

```
Collecting bayesian-optimization
  Downloading bayesian_optimization-1.4.3-py3-none-any.whl (18 kB)
Requirement already satisfied: numpy>=1.9.0 in /usr/local/lib/python3.10/dist-packages (from bayesian-optimization) (1.25.2)
Requirement already satisfied: scipy>=1.0.0 in /usr/local/lib/python3.10/dist-packages (from bayesian-optimization) (1.11.4)
Requirement already satisfied: scikit-learn>=0.18.0 in /usr/local/lib/python3.10/dist-packages (from bayesian-optimization) (1.2.2)
Collecting colorama>=0.4.6 (from bayesian-optimization)
  Downloading colorama-0.4.6-py2.py3-none-any.whl (25 kB)
Requirement already satisfied: joblib>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from scikit-learn>=0.18.0->bayesian-optimization) (1.4.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn>=0.18.0->bayesian-optimization) (3.2.0)
Installing collected packages: colorama, bayesian-optimization
Successfully installed bayesian-optimization-1.4.3 colorama-0.4.6
```


```
import numpy as np
import pandas as pd
from keras.datasets import mnist
from sklearn.model_selection import train_test_split, cross_val_score, GridSearchCV
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
```

✓ IMPORTING DATASET

```
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz
11490434/11490434 [=====] - 0s 0us/step
```

```
from matplotlib import pyplot
for i in range(3):
    pyplot.subplot(330 + 1 + i)
    pyplot.imshow(X_train[i], cmap=pyplot.get_cmap('gray'))
    pyplot.show()
```



```
X_train = X_train.reshape(-1, 28*28)
X_test = X_test.reshape(-1, 28*28)
```

Normalizing the pixel values to range [0, 1]

```
X_train = X_train / 255.0
X_test = X_test / 255.0
```

✓ IMPLEMENTING TRAINING MODELS

DECISION TREE CLASSIFIER

```
dt_classifier = DecisionTreeClassifier(random_state=42)
dt_scores = cross_val_score(dt_classifier, X_train, y_train, cv=5)
print("Decision Tree Classifier Accuracy: {:.2f}%".format(dt_scores.mean() * 100))
```

```
Decision Tree Classifier Accuracy: 86.65%
```

RANDOM FOREST CLASSIFIER

```
rf_classifier = RandomForestClassifier(random_state=42)
rf_scores = cross_val_score(rf_classifier, X_train, y_train, cv=5)
print("Random Forest Classifier Accuracy: {:.2f}%".format(rf_scores.mean() * 100))
```

```
Random Forest Classifier Accuracy: 96.64%
```

NAIVE BAYES CLASSIFIER

```
nb_classifier = GaussianNB()
nb_scores = cross_val_score(nb_classifier, X_train, y_train, cv=5)
print("Naïve Bayes Classifier Accuracy: {:.2f}%".format(nb_scores.mean() * 100))
```

Naïve Bayes Classifier Accuracy: 56.18%

KNN CLASSIFIER

```
knn_classifier = KNeighborsClassifier()
knn_scores = cross_val_score(knn_classifier, X_train, y_train, cv=5)
print("KNN Classifier Accuracy: {:.2f}%".format(knn_scores.mean() * 100))
```

KNN Classifier Accuracy: 96.93%

NEURAL NETWORK CLASSIFIER

```
nn_classifier = MLPClassifier(random_state=42)
nn_scores = cross_val_score(nn_classifier, X_train, y_train, cv=5)
print("Neural Network Classifier Accuracy: {:.2f}%".format(nn_scores.mean() * 100))
```

Neural Network Classifier Accuracy: 97.47%

Comparing through Graph

```
import matplotlib.pyplot as plt
import seaborn as sns

sns.set(style='whitegrid')
palette = sns.color_palette('viridis')

classifiers = ['Decision Tree', 'Random Forest', 'Naive Bayes', 'K-Nearest Neighbors', 'Multilayer Perceptron']
# classifier_accuracies=[86.65, 96.64, 56.18, 96.93, 97.47]
classifier_accuracies=[dt_scores.mean(), rf_scores.mean(), nb_scores.mean(), knn_scores.mean(), nn_scores.mean()]


plt.figure(figsize=(10, 6))
bars = plt.bar(classifiers, classifier_accuracies, color=palette, width=0.5)

plt.xlabel('Classifier', fontsize=12, fontweight='bold')
plt.ylabel('Accuracy', fontsize=12, fontweight='bold')
plt.title('Comparison of Classifier Accuracies on MNIST Dataset', fontsize=14, fontweight='bold')

plt.grid(axis='y', linestyle='--', alpha=0.6)

for bar in bars:
    yval = bar.get_height()
    plt.text(bar.get_x() + bar.get_width() / 2, yval + 0.01, f'{yval:.2f}', ha='center', va='bottom', fontsize=10, fontweight='bold')

plt.show()
```



```
data = {'Classifier': classifiers, 'Accuracy': classifier_accuracies}
df = pd.DataFrame(data)
df_sorted = df.sort_values(by='Accuracy', ascending=False)
df_sorted = df_sorted.reset_index(drop=True)
df_sorted['Rank'] = df_sorted.index + 1
print(df_sorted)

best_classifier = df_sorted.iloc[0]['Classifier']
best_accuracy = df_sorted.iloc[0]['Accuracy']
print(f"\nThe classifier with the highest accuracy ({best_accuracy:.4f}) is {best_classifier}, making it the best for classification.")
```

	Classifier	Accuracy	Rank
0	Multilayer Perceptron	97.47	1
1	K-Nearest Neighbors	96.93	2
2	Random Forest	96.64	3
3	Decision Tree	86.65	4
4	Naive Bayes	56.18	5

The classifier with the highest accuracy (97.4700) is Multilayer Perceptron, making it the best for classification.

✓ EVALUATION METRICS

```
classifiers = [dt_classifier, rf_classifier, nb_classifier, knn_classifier, nn_classifier]
classifier_names = ["Decision Tree", "Random Forest", "Naive Bayes", "KNN", "Neural Network"]
```

```
for clf, name in zip(classifiers, classifier_names):
    clf.fit(X_train, y_train)
    y_pred = clf.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)
    report = classification_report(y_test, y_pred)
    conf_matrix = confusion_matrix(y_test, y_pred)
    print(f"\n{name} Evaluation Metrics:")
    print(f"Accuracy: {accuracy}")
    print(f"Classification Report:\n{report}")
    print(f"Confusion matrix:\n{conf_matrix}")
    print('-'*150)
```

	0	0.96	0.99	0.98	980
	1	0.95	1.00	0.98	1135
	2	0.98	0.96	0.97	1032
	3	0.96	0.97	0.97	1010
	4	0.98	0.96	0.97	982
	5	0.97	0.97	0.97	892
	6	0.98	0.99	0.98	958
	7	0.96	0.96	0.96	1028
	8	0.99	0.94	0.96	974
	9	0.96	0.95	0.95	1009

accuracy			0.97	10000
macro avg	0.97	0.97	0.97	10000
weighted avg	0.97	0.97	0.97	10000

Confusion matrix:

[[974	1	1	0	0	1	2	1	0	0]
[[0	1133	2	0	0	0	0	0	0	0]
[[11	8	991	2	1	0	1	15	3	0]
[[0	3	3	976	1	13	1	6	3	4]
[[3	7	0	0	944	0	4	2	1	21]
[[5	0	0	12	2	862	4	1	2	4]
[[5	3	0	0	3	2	945	0	0	0]
[[0	22	4	0	3	0	0	988	0	11]
[[8	3	5	13	6	12	5	5	913	4]
[[5	7	3	9	7	3	1	10	2	962]]

Neural Network Evaluation Metrics:

Accuracy: 0.9782

Classification Report:

	precision	recall	f1-score	support
0	0.98	0.99	0.99	980
1	0.99	0.99	0.99	1135
2	0.98	0.98	0.98	1032
3	0.98	0.98	0.98	1010
4	0.97	0.97	0.97	982
5	0.98	0.98	0.98	892
6	0.98	0.98	0.98	958
7	0.98	0.98	0.98	1028
8	0.97	0.97	0.97	974
9	0.97	0.97	0.97	1009

accuracy			0.98	10000
macro avg	0.98	0.98	0.98	10000
weighted avg	0.98	0.98	0.98	10000

Confusion matrix:

[[972	0	1	0	1	0	2	1	3	0]
[[0	1125	3	0	0	0	2	1	4	0]
[[4	1	1007	2	1	0	4	5	7	1]
[[0	0	1	989	0	6	0	3	4	7]
[[3	0	3	1	956	0	4	4	2	9]
[[2	0	0	6	3	871	4	0	4	2]
[[4	2	2	1	4	5	938	0	2	0]
[[1	2	11	2	0	0	0	1003	3	6]
[[4	1	3	5	8	2	2	2	943	4]
[[3	2	0	5	9	4	1	4	3	978]]

✓ PARAMETER TUNING

✓ Through Grid search

Defining parameter grid for each classifier

```
dt_classifier = DecisionTreeClassifier(random_state=42)
rf_classifier = RandomForestClassifier(random_state=42)
knn_classifier = KNeighborsClassifier()
nn_classifier = MLPClassifier(random_state=42)

param_grid_dt = {
```

```

    'criterion': ['gini', 'entropy'],
    'splitter': ['best', 'random'],
    'max_depth': [10, 20, 30,],
    'min_samples_split': [2, 5,],
    'min_samples_leaf': [1, 2, 4]
}

param_grid_rf = {
    'n_estimators': [50, 100],
    'criterion': ['gini', 'entropy'],
    'max_depth': [10, 20, 30],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4]
}

param_grid_knn = {
    'n_neighbors': [3, 5, 7],
    'weights': ['uniform', 'distance'],
    'metric': ['euclidean', 'manhattan']
}

param_grid_nn = {
    'hidden_layer_sizes': [(50,), (100,)],
    'activation': ['logistic', 'relu'],
    'solver': ['adam', 'sgd'],
    'alpha': [0.0001, 0.001, 0.01]
}

```

Grid search for each classifier

```

grid_search_dt = GridSearchCV(dt_classifier, param_grid_dt, cv=5)
grid_search_rf = GridSearchCV(rf_classifier, param_grid_rf, cv=5)
grid_search_knn = GridSearchCV(knn_classifier, param_grid_knn, cv=5)
grid_search_nn = GridSearchCV(nn_classifier, param_grid_nn, cv=5)

```

Fit the grid search objects

```

grid_search_dt.fit(X_train, y_train)
grid_search_rf.fit(X_train, y_train)
grid_search_knn.fit(X_train, y_train)
grid_search_nn.fit(X_train, y_train)

```

Best parameters for each classifier

```

print("Best parameters for Decision Tree:", grid_search_dt.best_params_)
print("Best parameters for Random Forest:", grid_search_rf.best_params_)
print("Best parameters for KNN:", grid_search_knn.best_params_)
print("Best parameters for Neural Network:", grid_search_nn.best_params_)

```

```

➤ Best parameters for Decision Tree: {'max_depth': 20, 'min_samples_leaf': 1, 'min_samples_split': 2}

```

✓ Through Parameter Search

✓ Random Forest

Random Search

```

from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import randint, uniform

```

```

# import warnings filter
from warnings import simplefilter
# ignore all future warnings
simplefilter(action='ignore', category=FutureWarning)

```

```

param_dist = {
    'n_estimators': randint(10, 200),
    'max_depth': randint(1, 20),
    'min_samples_split': randint(2, 10),
    'min_samples_leaf': randint(1, 10),
    'max_features': ['auto', 'sqrt', 'log2']
}

```

```
rf_classifier = RandomForestClassifier(random_state=42)
```

```
random_search = RandomizedSearchCV(rf_classifier, param_distributions=param_dist, n_iter=20, cv=5, random_state=42)
```

```
random_search.fit(X_train, y_train);
```

```

print(f"Best parameters: {random_search.best_params_}")
print(f"Best cross-validation score: {random_search.best_score_:.4f}")

```

```

➤ Best parameters: {'max_depth': 1, 'max_features': 'auto', 'min_samples_leaf': 1, 'min_samples_split': 4, 'n_estimators': 10}
Best cross-validation score: 0.4660

```

Bayesian Optimization

```

from bayes_opt import BayesianOptimization

def rf_cv(n_estimators, max_depth, min_samples_split, min_samples_leaf):
    rf = RandomForestClassifier(
        n_estimators=int(n_estimators),
        max_depth=int(max_depth),
        min_samples_split=int(min_samples_split),
        min_samples_leaf=int(min_samples_leaf),
        random_state=42
    )
    cv_scores = cross_val_score(rf, X_train, y_train, cv=5)
    return cv_scores.mean()

pbounds = {
    'n_estimators': (10, 200),
    'max_depth': (1, 20),
    'min_samples_split': (2, 10),
    'min_samples_leaf': (1, 10),
}

# Initialize the Bayesian optimizer
optimizer = BayesianOptimization(
    f=rf_cv,
    pbounds=pbounds,
    random_state=42,
)

# Optimize
optimizer.maximize(init_points=5, n_iter=15)

# Print the best parameters and the corresponding best score
print(f"Best parameters: {optimizer.max['params']}")
print(f"Best cross-validation score: {optimizer.max['target']:.4f}")

```

Decision Tree

Random Search

```

param_dist_dt = {
    'max_depth': randint(1, 20),
    'min_samples_split': randint(2, 10),
    'min_samples_leaf': randint(1, 10)
}

dt_classifier = DecisionTreeClassifier(random_state=42)

random_search_dt = RandomizedSearchCV(
    dt_classifier, param_distributions=param_dist_dt, n_iter=1, cv=5, random_state=42
)
random_search_dt.fit(X_train, y_train)

print(f"Decision Tree Random Search - Best Parameters: {random_search_dt.best_params_}")
print(f"Decision Tree Random Search - Best CV Score: {random_search_dt.best_score_.4f}")

```

Decision Tree Random Search - Best Parameters: {'max_depth': 7, 'min_samples_leaf': 4, 'min_samples_split': 6}
 Decision Tree Random Search - Best CV Score: 0.7746

Bayesian Optimization

```

# Define function to optimize
def dt_cv(max_depth, min_samples_split, min_samples_leaf):
    dt = DecisionTreeClassifier(
        max_depth=int(max_depth),
        min_samples_split=int(min_samples_split),
        min_samples_leaf=int(min_samples_leaf),
        random_state=42
    )
    cv_scores = cross_val_score(dt, X_train, y_train, cv=5)
    return cv_scores.mean()

# Define parameter bounds for Bayesian optimization
pbounds_dt = {
    'max_depth': (1, 20),
    'min_samples_split': (2, 10),
    'min_samples_leaf': (1, 10)
}

# Initialize Bayesian optimizer
optimizer_dt = BayesianOptimization(
    f=dt_cv,
    pbounds=pbounds_dt,
    random_state=42
)

# Optimize
optimizer_dt.maximize(init_points=5, n_iter=15)

# Print the best parameters and score
print(f"Decision Tree Bayesian Optimization - Best Parameters: {optimizer_dt.max['params']}")
print(f"Decision Tree Bayesian Optimization - Best CV Score: {optimizer_dt.max['target']:.4f}")

```

iter	target	max_depth	min_sa...	min_sa...
0	0.7746	7	4	6

1	0.8051	8.116	9.556	7.856	
2	0.8667	12.37	2.404	3.248	
3	0.3404	2.104	8.796	6.809	
4	0.869	14.45	1.185	9.759	
5	0.8685	16.82	2.911	3.455	
6	0.866	14.05	6.92	7.359	

Decision Tree Bayesian Optimization - Best Parameters: {'max_depth': 14.453378978124864, 'min_samples_leaf': 1.185260448662222, 'min_samples_split': 10}
Decision Tree Bayesian Optimization - Best CV Score: 0.8690

▼ KNN

Random Search

```
# Define parameter distributions for KNN
param_dist_knn = {
    'n_neighbors': randint(1, 20),
    'weights': ['uniform', 'distance'],
    'algorithm': ['auto', 'ball_tree', 'kd_tree', 'brute'],
    'leaf_size': randint(10, 50),
    'p': randint(1, 2)
}

# Initialize K-Nearest Neighbors Classifier
knn_classifier = KNeighborsClassifier()

# Perform Random Search
random_search_knn = RandomizedSearchCV(
    knn_classifier, param_distributions=param_dist_knn, n_iter=20, cv=5, random_state=42
)
random_search_knn.fit(X_train, y_train)

# Print the best parameters and score
print(f"K-Nearest Neighbors Random Search - Best Parameters: {random_search_knn.best_params_}")
print(f"K-Nearest Neighbors Random Search - Best CV Score: {random_search_knn.best_score_:.4f}")
```

Bayesian Optimization

```
# Define function to optimize
def knn_cv(n_neighbors, leaf_size, p):
    knn = KNeighborsClassifier(
        n_neighbors=int(n_neighbors),
        leaf_size=int(leaf_size),
        p=int(p),
        weights='uniform',
        algorithm='auto'
    )
    cv_scores = cross_val_score(knn, X_train, y_train, cv=5)
    return cv_scores.mean()

# Define parameter bounds for Bayesian optimization
pbounds_knn = {
    'n_neighbors': (1, 20),
    'leaf_size': (10, 50),
    'p': (1, 2)
}

# Initialize Bayesian optimizer
optimizer_knn = BayesianOptimization(
    f=knn_cv,
    pbounds=pbounds_knn,
    random_state=42
)

# Optimize
optimizer_knn.maximize(init_points=5, n_iter=15)

# Print the best parameters and score
print(f"K-Nearest Neighbors Bayesian Optimization - Best Parameters: {optimizer_knn.max['params']}")
print(f"K-Nearest Neighbors Bayesian Optimization - Best CV Score: {optimizer_knn.max['target']:.4f}")
```

▼ Neural Networks

Random Search

```
# Define parameter distributions for MLP
param_dist_mlp = {
    'hidden_layer_sizes': [(50,), (100,), (50, 50), (100, 100)],
    'activation': ['relu', 'tanh', 'logistic'],
    'solver': ['adam', 'sgd'],
    'alpha': uniform(0.0001, 0.1),
    'learning_rate': ['constant', 'invscaling', 'adaptive'],
    'learning_rate_init': uniform(0.001, 0.1),
    'max_iter': randint(100, 1000)
}

# Initialize MLPClassifier
mlp_classifier = MLPClassifier(random_state=42)

# Perform Random Search
random_search_mlp = RandomizedSearchCV(
```

```
    mlp_classifier, param_distributions=param_dist_mlp, n_iter=20, cv=5, random_state=42
)
random_search_mlp.fit(X_train, y_train)

# Print the best parameters and score
print(f"MLP Random Search - Best Parameters: {random_search_mlp.best_params}")
print(f"MLP Random Search - Best CV Score: {random_search_mlp.best_score_:.4f}")
```

Bayesian Optimization

```
# Define function to optimize
def mlp_cv(hidden_layer_sizes, activation, solver, alpha, learning_rate, learning_rate_init, max_iter):
    mlp = MLPClassifier(
        hidden_layer_sizes=eval(hidden_layer_sizes),
        activation=activation,
        solver=solver,
        alpha=alpha,
        learning_rate=learning_rate,
        learning_rate_init=learning_rate_init,
        max_iter=max_iter,
        random_state=42
    )
    cv_scores = cross_val_score(mlp, X_train, y_train, cv=5)
    return cv_scores.mean()

# Define parameter bounds for Bayesian optimization
pbounds_mlp = {
    'hidden_layer_sizes': ['(50,)', '(100,)', '(50,50)', '(100,100)'],
    'activation': ['relu', 'tanh', 'logistic'],
    'solver': ['adam', 'sgd'],
    'alpha': (0.0001, 0.1),
    'learning_rate': ['constant', 'invscaling', 'adaptive'],
    'learning_rate_init': (0.001, 0.1),
    'max_iter': (100, 1000)
}

# Initialize Bayesian optimizer
optimizer_mlp = BayesianOptimization(
    f=mlp_cv,
    pbounds=pbounds_mlp,
    random_state=42
)

# Optimize
optimizer_mlp.maximize(init_points=5, n_iter=15)

# Print the best parameters and score
print(f"MLP Bayesian Optimization - Best Parameters: {optimizer_mlp.max['params']}")
print(f"MLP Bayesian Optimization - Best CV Score: {optimizer_mlp.max['target']:.4f}")
```